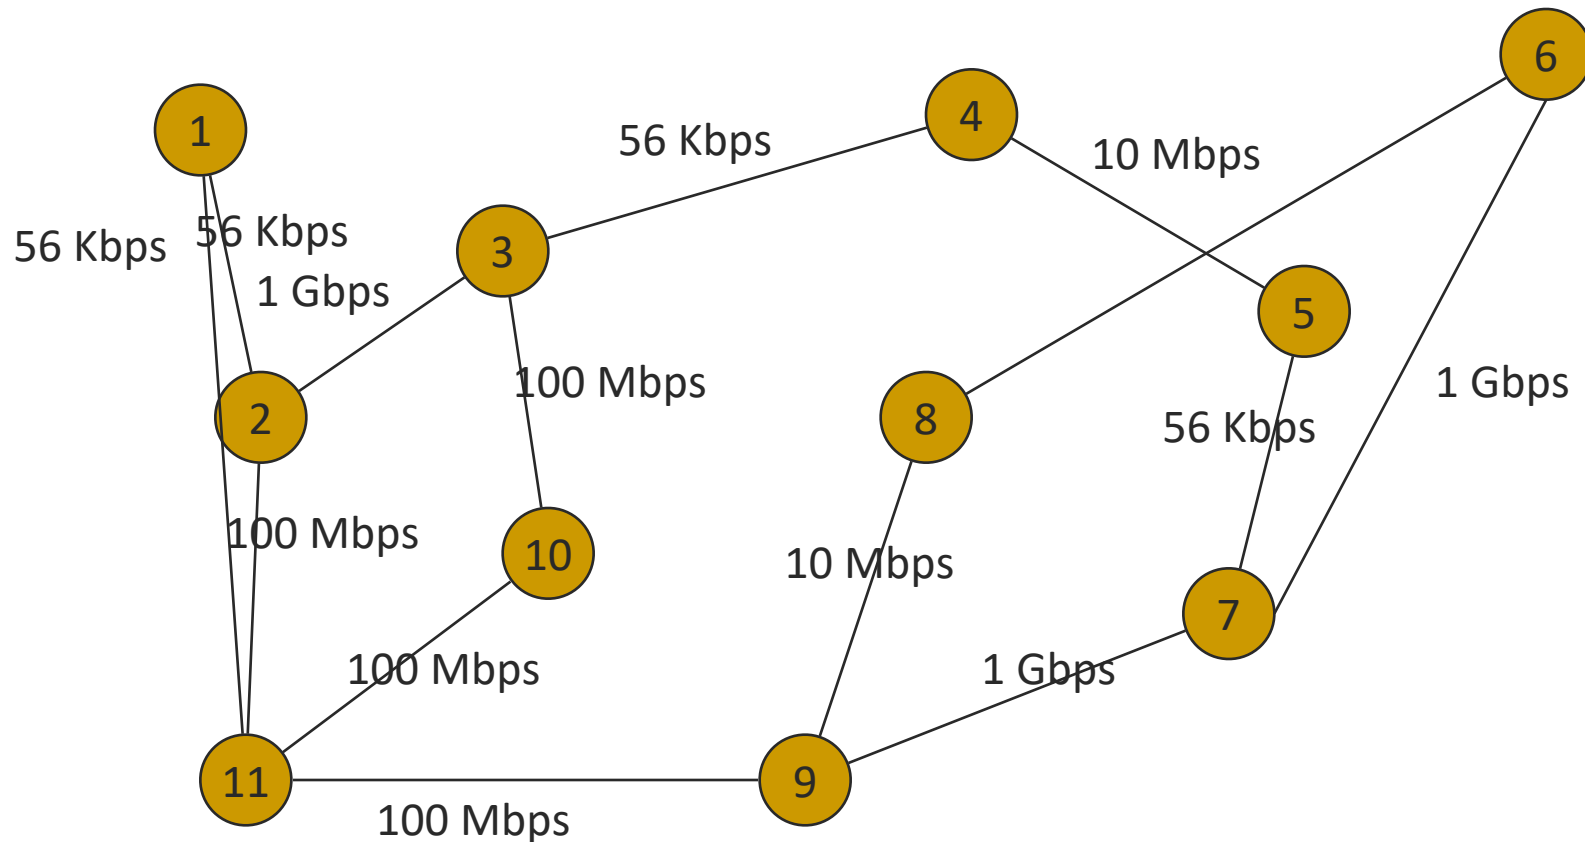


Shortest Paths

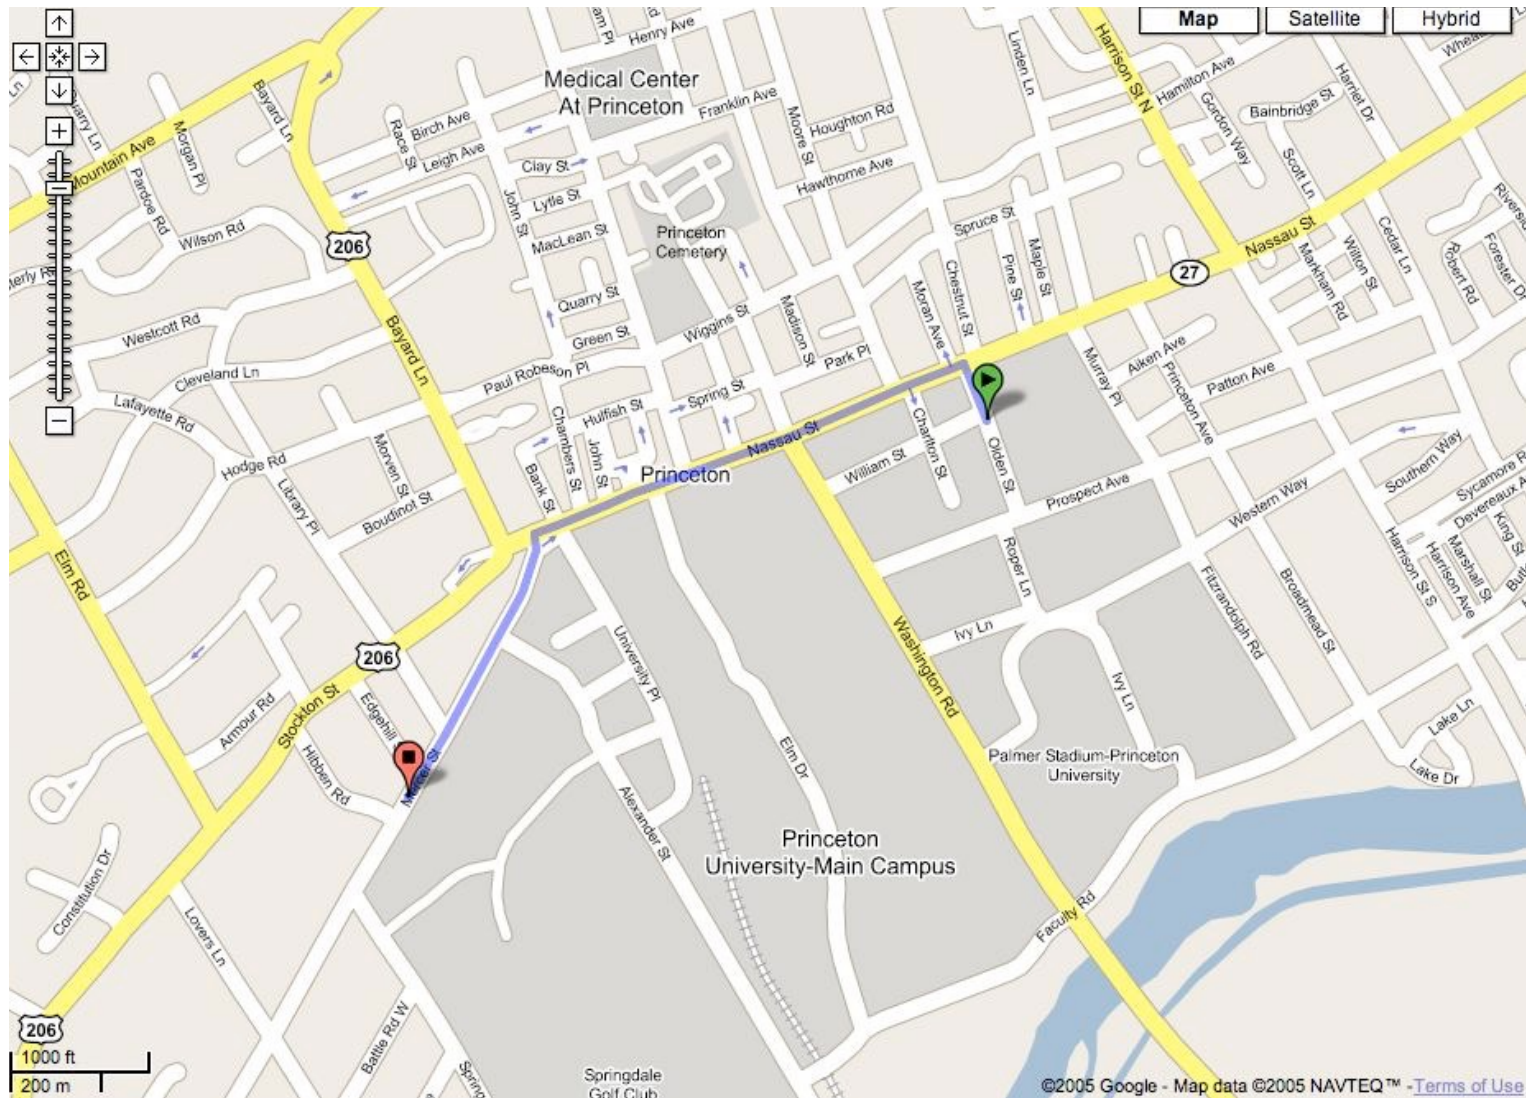
in edge-weighted graphs

Communication Speeds in a Computer Network

Find fastest way to route a data packet between two computers



Google maps



Shortest path applications

- PERT/CPM.
- Map routing.
- **Seam carving.**
- Texture mapping.
- Robot navigation.
- Typesetting in TeX.
- Urban traffic planning.
- Optimal pipelining of VLSI chip.
- Telemarketer operator scheduling.
- Routing of telecommunications messages.
- Network routing protocols (OSPF, BGP, RIP).
- Exploiting **arbitrage** opportunities in currency exchange.
- Optimal truck routing through given traffic congestion pattern.



http://en.wikipedia.org/wiki/Seam_carving



Reference: Network Flows: Theory, Algorithms, and Applications, R. K. Ahuja, T. L. Magnanti, and J. B. Orlin, Prentice Hall, 1993.

Shortest Path problems

- Find a shortest path between two given vertices
- Single source shortest paths
- Single sink shortest paths
- All pair shortest paths

Single Source Shortest Path problems

- Undirected graphs with non-negative edge weights
- Directed graphs with non-negative edge weights
- Directed graphs with arbitrary weights

Single Source Shortest Paths

- If graph is not weighted (all edge-weights are unit-weight): BFS works
- Now assume: graph is edge-weighted
 - Every edge is associated with a positive number
 - Possible weights: integers, real numbers, rational numbers
 - Edge-weights can represent: distance, connection cost, affinity

Single Source Shortest Paths and shortest distances

- **Input:** An edge-weighted undirected graph and a source node v with: for every edge e edge-weight $w(e) > 0$
- **Output:** All single-source shortest paths (and their weight) for v in G : for every node $w \neq v$ in G a shortest path from v to w .
 - Here, a *path* p from v to w consisting of edges e_0, e_1, \dots, e_{k-1} is shortest in G , if its length

$$w(p) = \sum_{i=0}^{k-1} w(e_i)$$

is minimum (i.e., there is no path from v to w in G that is shorter).

Let $d(u)$ = length of shortest path from v to u

Algorithm

DijkstraShortestPaths(G, v)

Input: A simple undirected graph G with nonnegative edge-weights, a distinguished vertex v in G

Output: A label $D[u]$ for each vertex u in G such that $D[u] = d(u)$

Algorithm

DijkstraShortestPaths(G, v)

$D[v] \leftarrow 0$

for each vertex $u \neq v$ of G **do**

$D[u] \leftarrow +\infty$

Let Q be a priority queue containing all vertices of G using $D[.]$ as keys

while Q is not empty **do**

$u \leftarrow Q.\text{removeMin}()$ // u is added to T

for each vertex $z \in N(u)$ with $z \in Q$ **do**

if $D[u] + w((u, z)) < D[z]$ **then**


Relaxation

$D[z] \leftarrow D[u] + w((u, z))$


update z 's key in Q to $D[z]$

return D


Dijkstra's algorithm: a greedy algorithm




Dijkstra's algorithm: Initializing




Dijkstra's algorithm: Initializing Cloud C (consisting of "solved" subgraph)




Dijkstra's algorithm: pull v into C




Dijkstra's algorithm: update C 's neighborhood




Dijkstra's algorithm: pick closest vertex u outside C




Dijkstra's algorithm: pull u into C




Dijkstra's algorithm: update C 's neighborhood




Dijkstra's algorithm: pick closest vertex u outside C




Dijkstra's algorithm: pull u into C




Dijkstra's algorithm: update C 's neighborhood




Dijkstra's algorithm: pick closest vertex u outside C




Dijkstra's algorithm: pull u into C




Dijkstra's algorithm: update C 's neighborhood




Dijkstra's algorithm: pick closest vertex u outside C




Dijkstra's algorithm: pull u into C




Dijkstra's algorithm: update C 's neighborhood




Dijkstra's algorithm: pick closest vertex u outside C




Dijkstra's algorithm: pull u into C




Dijkstra's algorithm: update C 's neighborhood




Dijkstra's algorithm: pick closest vertex u outside C




Dijkstra's algorithm: pull u into C




Dijkstra's algorithm: update C 's neighborhood




Dijkstra's algorithm: pick closest vertex u outside C




Dijkstra's algorithm: pull u into C




Dijkstra's algorithm: update C 's neighborhood




Dijkstra's algorithm: pick closest vertex u outside C




Dijkstra's algorithm: pull u into C




Dijkstra's algorithm: update C 's neighborhood



Dijkstra's algorithm: pick closest vertex u outside C



Dijkstra's algorithm: pull u into C



Algorithm

DijkstraShortestPaths(G, v)

$D[v] \leftarrow 0$

for each vertex $u \neq v$ of G **do**

$D[u] \leftarrow +\infty$

Let Q be a priority queue containing all vertices of G using $D[.]$ as keys

while Q is not empty **do**

$u \leftarrow Q.\text{removeMin}()$ // u is added to T

for each vertex $z \in N(u)$ with $z \in Q$ **do**

if $D[u] + w((u, z)) < D[z]$ **then**

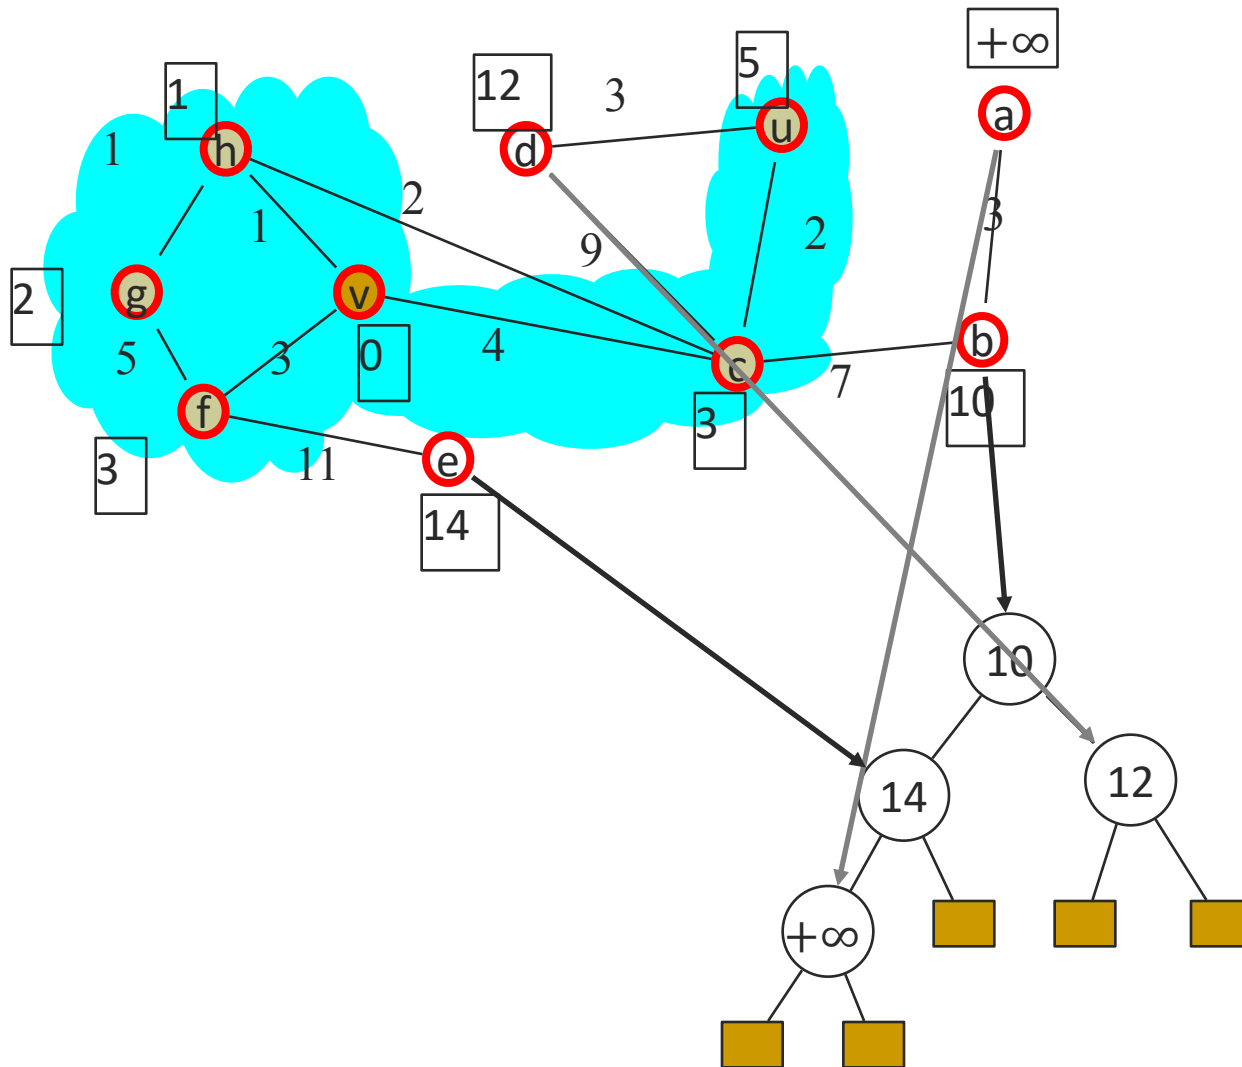
Relaxation

$D[z] \leftarrow D[u] + w((u, z))$

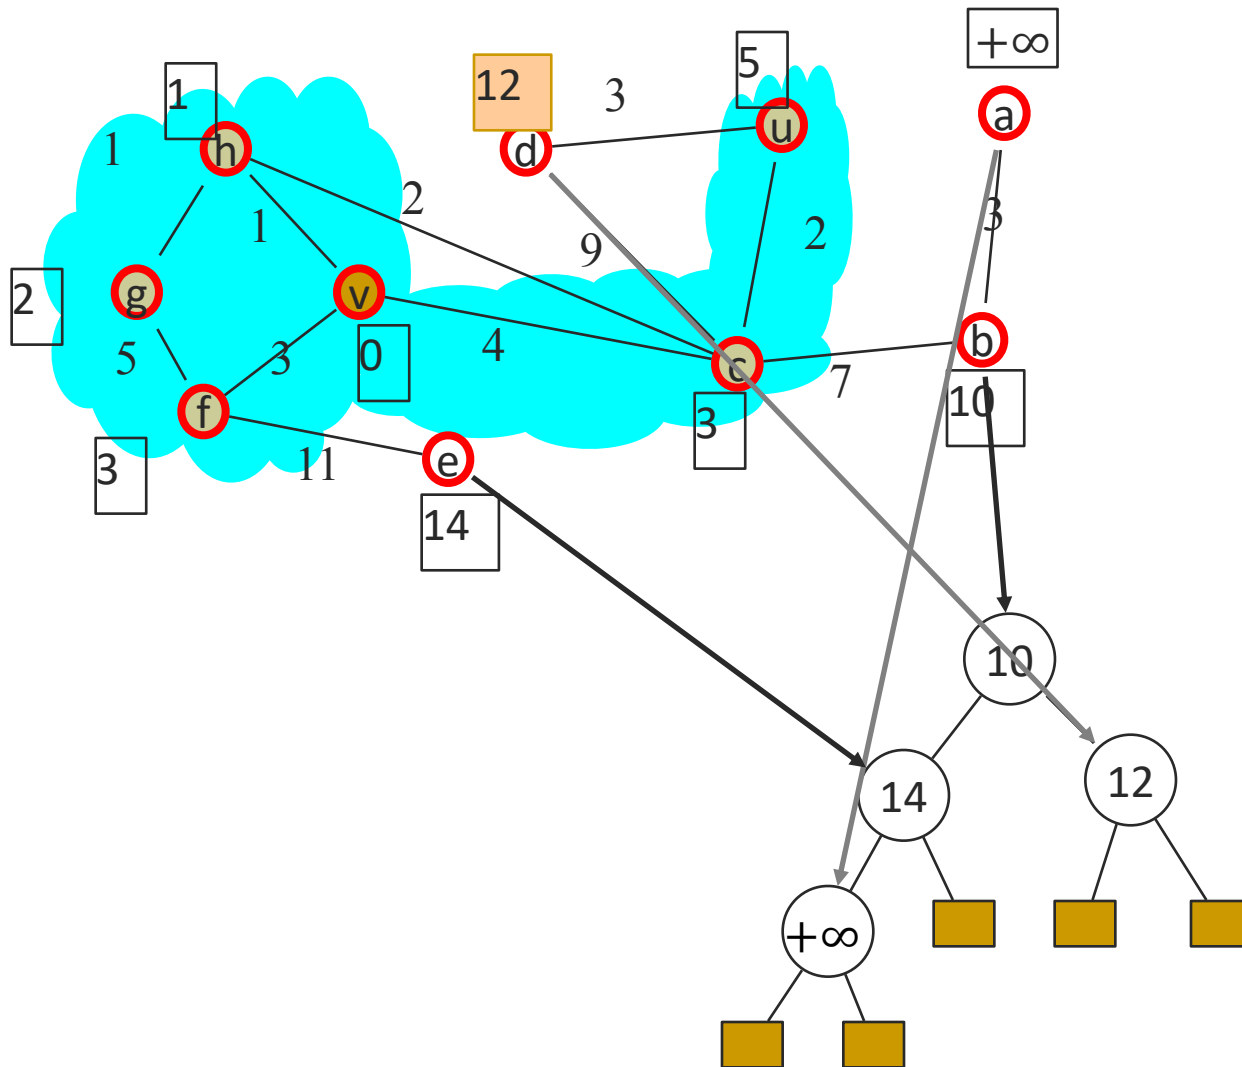
update z 's key in Q to $D[z]$

return D

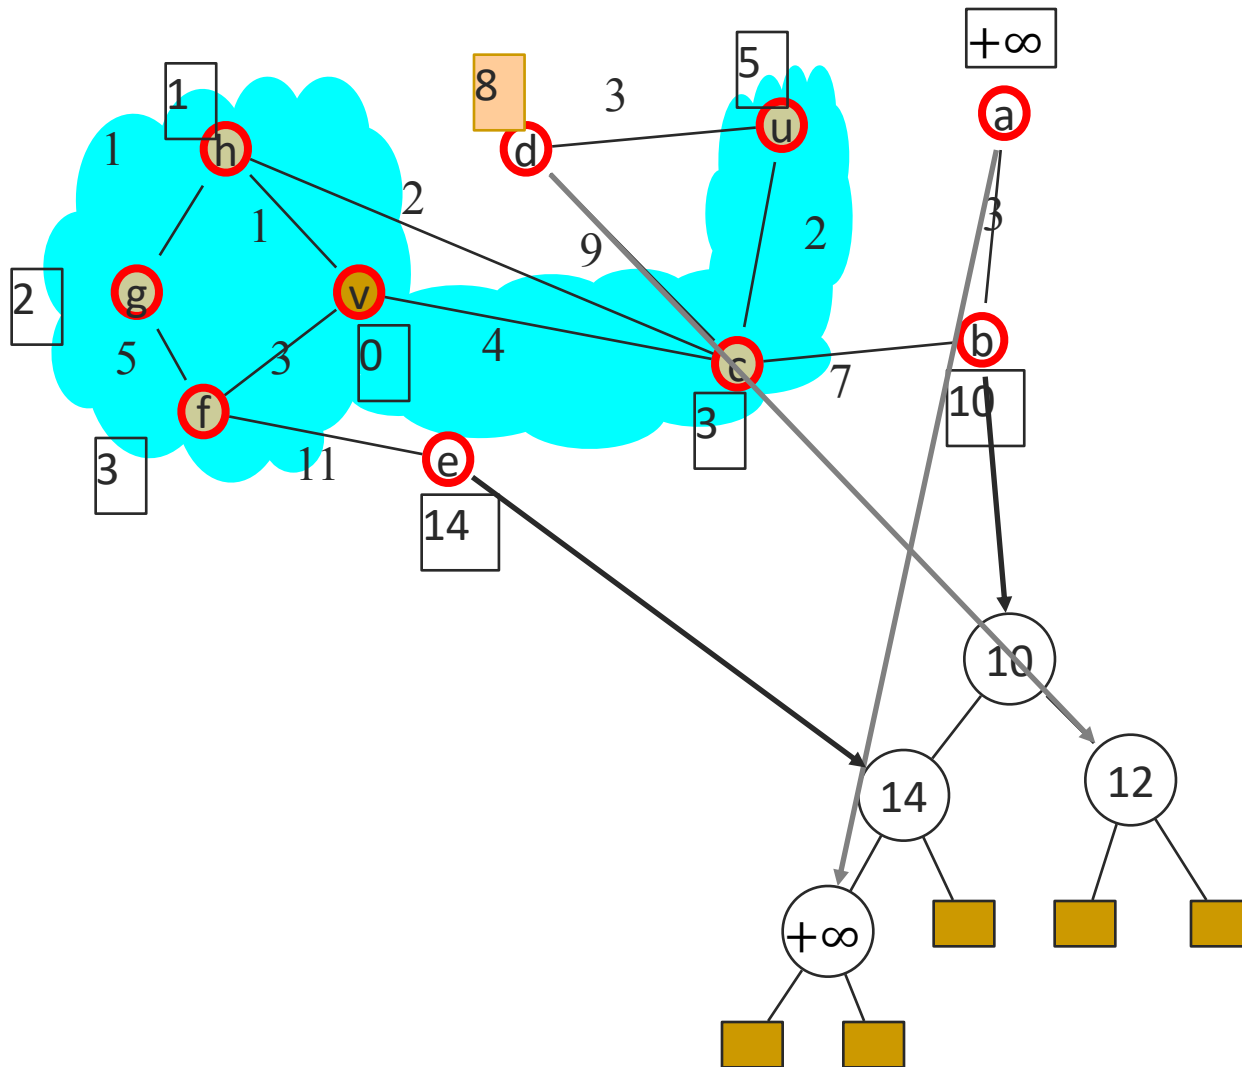
Dijkstra's algorithm using heaps



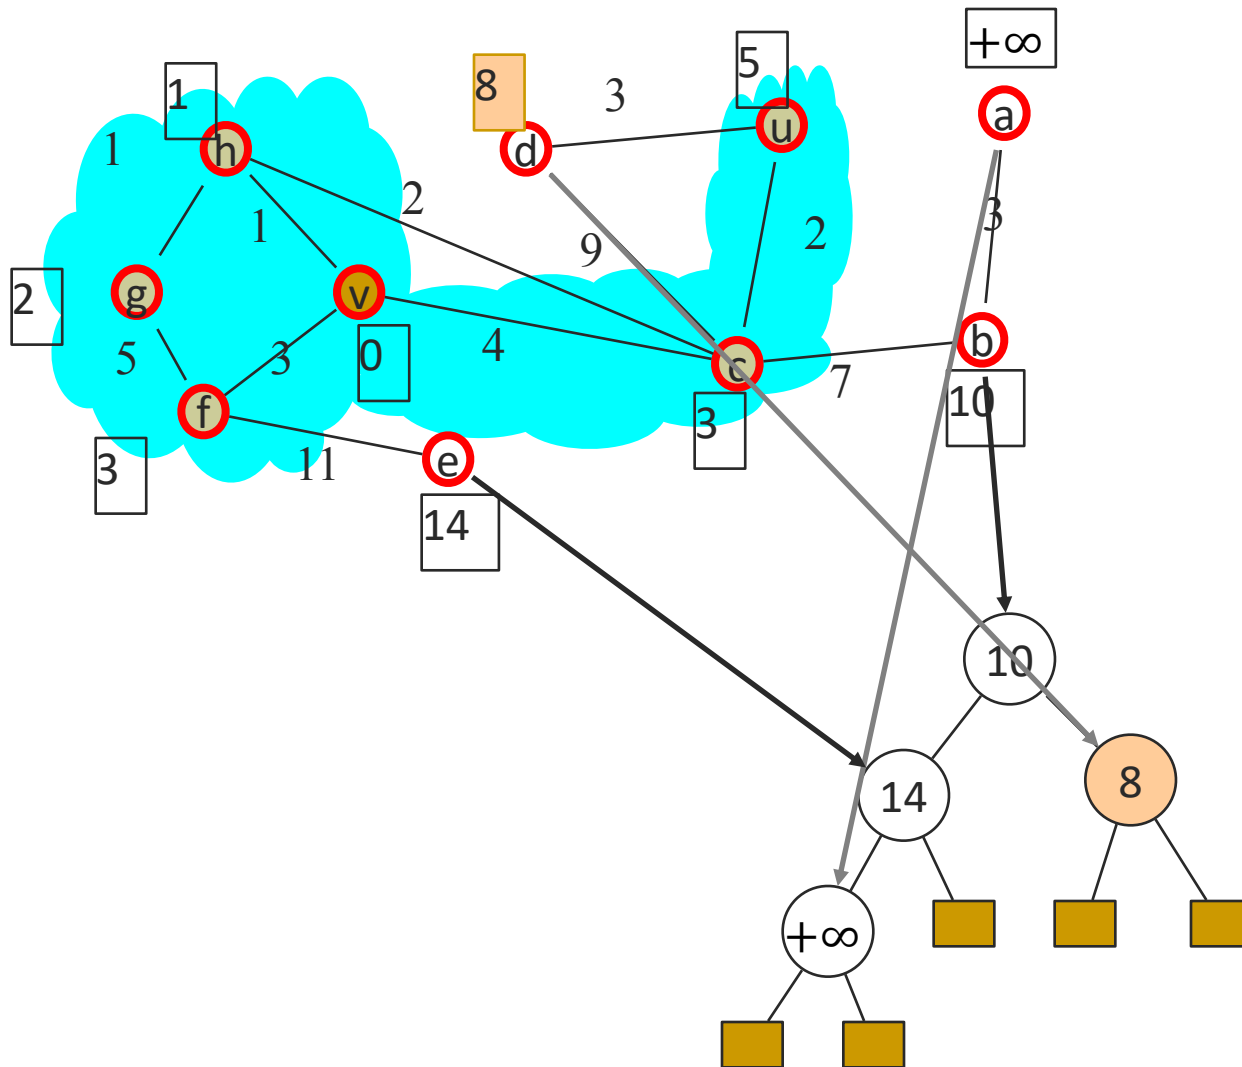
Dijkstra's algorithm using heaps



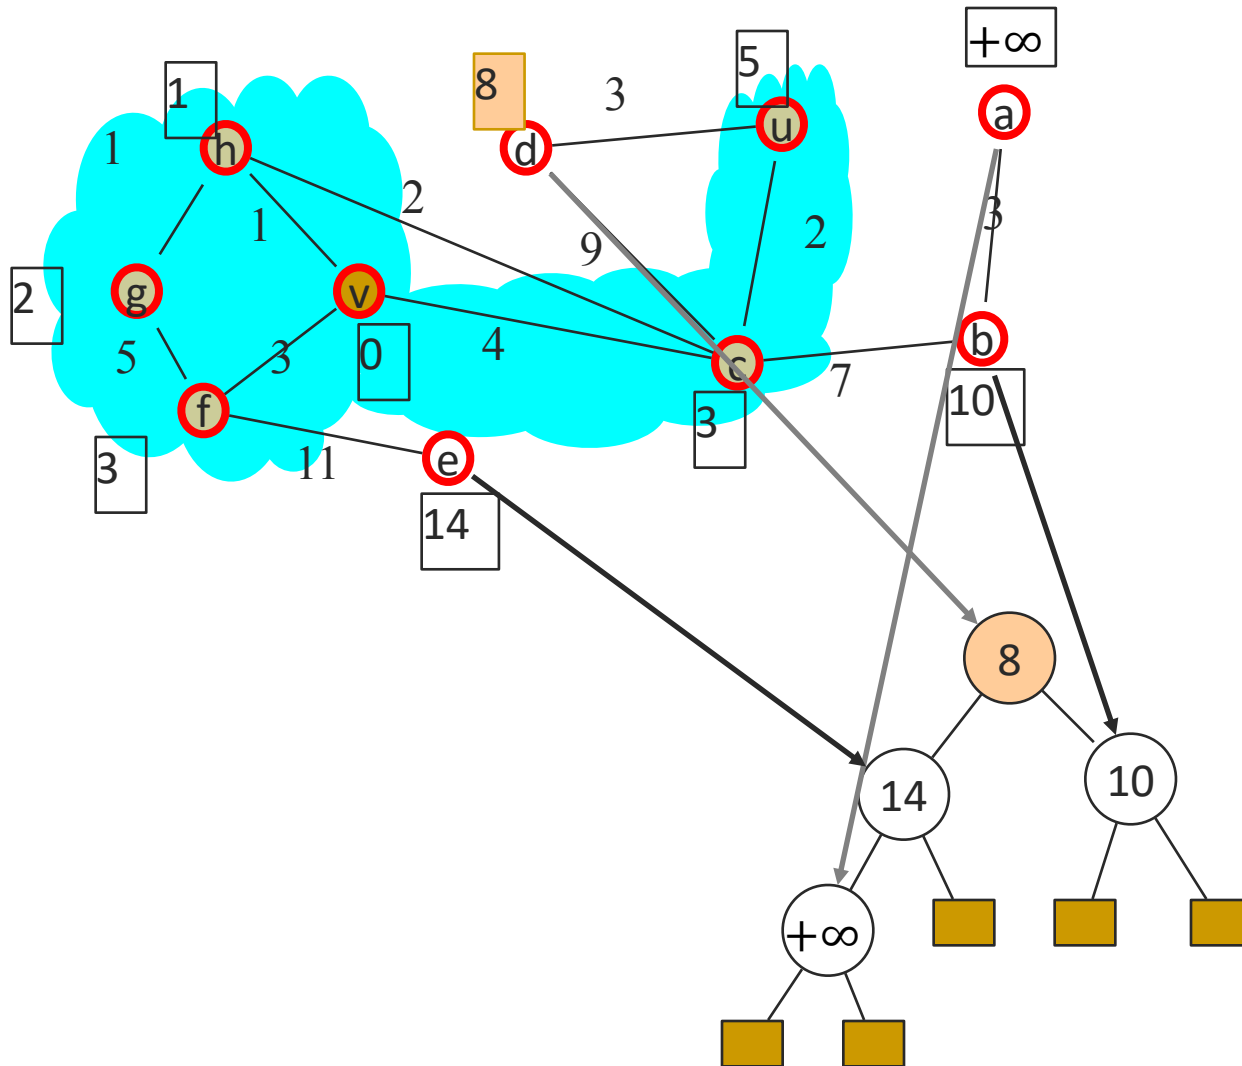
Dijkstra's algorithm using heaps



Dijkstra's algorithm using heaps



Dijkstra's algorithm using heaps



Running time

$D[v] \leftarrow 0$

for each vertex $u \neq v$ of G **do**

$D[u] \leftarrow +\infty$

Let Q be a priority queue containing all vertices of G using $D[.]$ as keys

while Q is not empty **do**

$u \leftarrow Q.\text{removeMin}()$ // u is added to cloud

for each vertex $z \in N(u)$ with $z \in Q$ **do**

if $D[u] + w((u, z)) < D[z]$ **then**

Relaxation

$D[z] \leftarrow D[u] + w((u, z))$

update z 's key in Q to $D[z]$

return D

Running time for $G=(V,E)$ with $|V|=n$ and $|E|=m$

- Insertion of vertices in priority queue Q
 - $O(n)$ when using bottom-up heap construction
- While loop:
 - Per iteration:
 - Remove vertex from Q $O(\log n)$
 - Relaxation $O(\deg(u) \log(n))$
 - $\sum_{u \in G} (1 + \deg(u)) \log n$ is $O((n + m) \log n)$
- Overall running time: $O(m \log n)$

In real life applications

- Often the graphs are sparse
- Then $O(m \log n)$ may be $O(n \log n)$

Algorithm

Correctness

DijkstraShortestPaths(G, v)

$D[v] \leftarrow 0$

for each vertex $u \neq v$ of G **do**

$D[u] \leftarrow +\infty$

Let Q be a priority queue containing all vertices of G using $D[.]$ as keys

while Q is not empty **do**

$u \leftarrow Q.\text{removeMin}()$ // u is added to T

for each vertex $z \in \text{adj}(u)$ with $z \in Q$ **do**

if $D[u] + w((u, z)) < D[z]$ **then**

Relaxation

$D[z] \leftarrow D[u] + w((u, z))$

update z 's key in Q to $D[z]$

return D

Invariants

Let $d(u)$ =distance of u from v

1. For each node u in T , $D[u] = d(u)$
2. For each node u not in T , $D[u]$ = length of shortest path from v to u without the use of other nodes outside of T

$+\infty$ denotes that the node cannot be reached yet from v via T nodes only

→ For all nodes u in T , $D[u] \geq d(u)$

Proof that invariants hold:

Initially they hold:

$$D(v)=0, D(u)=\infty$$

Consider the **first time** they fail:

CASE A: If Invariant 1 holds, i.e, $D(u)=d(u)$ for all u in T , but Invariant 2 fails for some z outside T when we add u to T , i.e, $D(z)$ fails to be the length of the shortest path from v to z passing only through nodes in T .

Claim: This can only happen if there is a shorter path of length $d' < D(z)$ from v to z whose last edge is $\{u, z\}$:


Why?

Claim: Assuming Invariant 1 holds, Invariant 2 is violated only if there is a shorter path from v to z whose last edge is $\{u, z\}$: Why?

Proof of Claim: Suppose to the contrary, the last edge of the shorter path is $\{x, z\}$, $x \neq u$, x in T .


We would have previously set $D(z)$ to $D(x) + w(x, z)$ when we relaxed x .

By Invariant 1, $D(x) = d(x)$ since x is in T , so $D(z)$ would have been set to $d(x) + w(x, z) = d' < D(z)$ giving a contradiction.



Since u is the last node in the shortest path to z passing through only T , then the relaxation step correctly sets

- $D(z) = d(v, u) + w((u, z))$ so Invariant 2 is NOT violated after the relaxation with respect to u is completed.



Proof that invariants hold (cont'd):

CASE B: (Invariant 1 fails first)

Consider the first time Invariant 1 fails:

u is added to T , $D(u)$ is the length of the shortest path from u to v through T but $D(u) \neq d(u)$

Let P = shortest path from v to u

- Note: For all u , $D(u) \geq d(u)$
- Let y = last node in T on path P
- By Invariant 1, $D[y] = d(y)$


There must be some $z \neq u$ on the path from y to u

Else $D(u) = d(u)$ (Why?)

- By Invariant 2, $D(z) \leq d(y) + w((y,z)) < d(u) < D(u)$

→ $D(z) < D(u)$

→ u would not be
added to T
before z



We've just show both invariants always hold

Let $d(u)$ =distance of u from v

1. For each node u in T , $D[u] = d(u)$
 2. For each node u not in T , $D[u]$ = length of shortest path from v to u without the use of other nodes outside of T
 $+\infty$ denotes that the node cannot be reached yet from v via T nodes only
- When the algorithm is finished, all nodes are in T and all $D[u]=d(u)$